

Optimizing Linear Algebra Performance for Programmers

Elliot Fiske

April 5, 2016

1 Multiplying Many Matrices and One Vector

Consider the following matrix operation:

$$Result = \underbrace{\mathbf{A} * \mathbf{B} * \mathbf{C} * \mathbf{D} * \mathbf{E}}_{100 \text{ by } 100} * \underbrace{\vec{x}}_{100 \text{ by } 1} \quad (1)$$

If this operation was done programmatically, with a compiler that evaluated each operation left-to-right, there would be over 4 million multiplication operations:

$$\underbrace{100 * 100 * 100 * 4}_{\substack{4 \text{ instances of multiply-} \\ \text{ing a } 100 \times 100 \text{ matrix} \\ \text{by a } 100 \times 100 \text{ matrix}}} + \underbrace{100 * 100}_{\substack{1 \text{ instance of} \\ \text{multiplying a} \\ 100 \times 100 \text{ ma-} \\ \text{trix by a } 100 \times 1 \\ \text{vector}}} = 4,010,000 \quad (2)$$

However, if we instead first evaluate $\mathbf{E} * \vec{x}$, then evaluate $\mathbf{D} * (\mathbf{E} * \vec{x})$, etc. there will be just over 1 million multiplication operations:

$$\underbrace{100 * 100 * 4}_{\substack{4 \text{ instances of multiply-} \\ \text{ing a } 100 \times 100 \text{ matrix} \\ \text{by a } 100 \times 1 \text{ vector}}} + \underbrace{100 * 100 * 100}_{\substack{1 \text{ instance of} \\ \text{multiplying a} \\ 100 \times 100 \text{ matrix} \\ \text{by a } 100 \times 100 \\ \text{matrix}}} = 1,040,000 \quad (3)$$

I used Eigen and C++ to practically test this, timing how long it took to do the approach from Equation 2 vs. the approach from Equation 3 with various matrix sizes.

| Matrix Dimension | Slow Method Time (μs) | Fast Method Time (μs) |
|------------------|------------------------------|------------------------------|
| 10 | 162 | 52 |
| 50 | 13,291 | 738 |
| 100 | 32,515 | 2561 |
| 300 | 537,397 | 17,399 |
| 1000 | 19,730,814 | 173,000 |

Table 1: Timing the dumb way vs. the smart way

For each dimension, the 'smart' method is orders of magnitude faster, becoming better as the dimension increases.

2 Calculating the Inverse of a Matrix

Consider the following matrix operation:

$$Result = \vec{x}^T * \mathbf{A}^T * \mathbf{B}^{-1} * \mathbf{A} * \vec{x} \quad (4)$$

One approach is to simply calculate the inverse of \mathbf{B} then multiply it with the rest of the terms. However, calculating the inverse of a matrix is a very expensive operation. Instead, consider rearranging the equation like so:

$$\begin{aligned} \mathbf{B}b &= \mathbf{A} * \vec{x} \\ Result &= \vec{x}^T * \mathbf{A}^T * b \end{aligned} \quad (5)$$

It is a much cheaper operation to solve a $\mathbf{A}\vec{x} = \mathbf{B}$ matrix problem than it is to find an inverse of a matrix. I used Eigen and C++ to practically test this, timing how long it takes to do the approach from Equation 4 vs. the approach from Equation 5.

Below a 50x50 matrix, the two methods are fairly equal in running time. However, at dimension 100 and above the solve() method is clearly superior.

| Matrix Dimension | <i>inverse()</i> Time (μ s) | <i>solve()</i> Time (μ s) |
|------------------|----------------------------------|--------------------------------|
| 10 | 61 | 70 |
| 50 | 278 | 267 |
| 100 | 670 | 935 |
| 300 | 9599 | 17,104 |
| 1000 | 320,196 | 244,522 |

Table 2: Using *inverse()* vs. using *solve()*